

# *High Performance Parallel Computing (1.)*

*overview, basics*

Imre Szeberényi, Éva Geist

BME IIT

<szebi@iit.bme.hu>



MŰEGYETEM 1782

# Course information

- Course site
  - <https://www.iit.bme.hu/targyak/BMEVIIIMA06>
  - teams + edu.vik.bme.hu
  - 2+1+0v (2 lecture + 1 practice + 0 lab + exam)
- Schedule
  - Lecture: Monday, 10:15-11:45, IL408
  - Practice: Thursday: 10:15-11:45, IB408 (biweekly, first on Sept. 14.)
- Assessment
  - Should not miss more than 30% of the classes
  - One final test
  - Homework – has to be completed by the exam
  - Written exam and homework presentation
  - 5 small knowledge test (optional for offering)

# Course information #2

## Grading:

IF best 3 from the 5 small tests are better than 80% (each) and the midterm test is better than 80%

- offered mark is:  $2 * \text{midterm test} + 1 * \text{HW}$

ELSE IF the midterm test is better than 40%

- mark is:  $2 * \text{written part of the exam} + 1 * \text{HW}$

ELSE

- mark is "not absolved,,

## HW

A parallel implementation and performance study of a problem or algorithm you are interested in. Implementations can use OpenMP, Pthreads, MPI, Charm, OpenCL, CUDA, ...

HW has to be demonstrated in person.

# Course information #3

## Usefull sources

Parhami, B., *Introduction to Parallel Processing: Algorithms and Architectures*, Plenum Press, 1999

Rauber, T. and G. Runger, *Parallel Programming for Multicore and Cluster Systems*, 2nd ed., Springer, 2013

Lawrence Livermore National Laboratory: Introduction to Parallel Computing Tutorial

<https://hpc.llnl.gov/training/tutorials/introduction-parallel-computing-tutorial>

Free on-line book (Creative Commons License)

Matloff, N., *Programming on Parallel Machines: GPU, Multicore, Clusters and More*, 341 pp., PDF file

<http://heather.cs.ucdavis.edu/~matloff/158/PLN/ParProcBook.pdf>

Useful free on-line course, sponsored by NVIDIA

“Introduction to Parallel Programming,” CPU/GPU-CUDA

<https://developer.nvidia.com/udacity-cs344-intro-parallel-programming>

# *Short History*

---

- In early 1950s Daniel Slotnick proposed a parallel machine, but John von Neumann dismissed because it requires „too many tubes”
- 1967: ILLIAC IV (256 proc, 200MFlop)
- Thinking Machines CM-1, CM-2 (1980)
- Cray-1

# *What is Parallel Processing?*



Parallelism = Concurrency

Doing more than one thing at a time

Has been around for decades, since early computers

I/O channels, DMA, device controllers, multiple ALUs

Multiple CPUs, Multiple Cores, ...

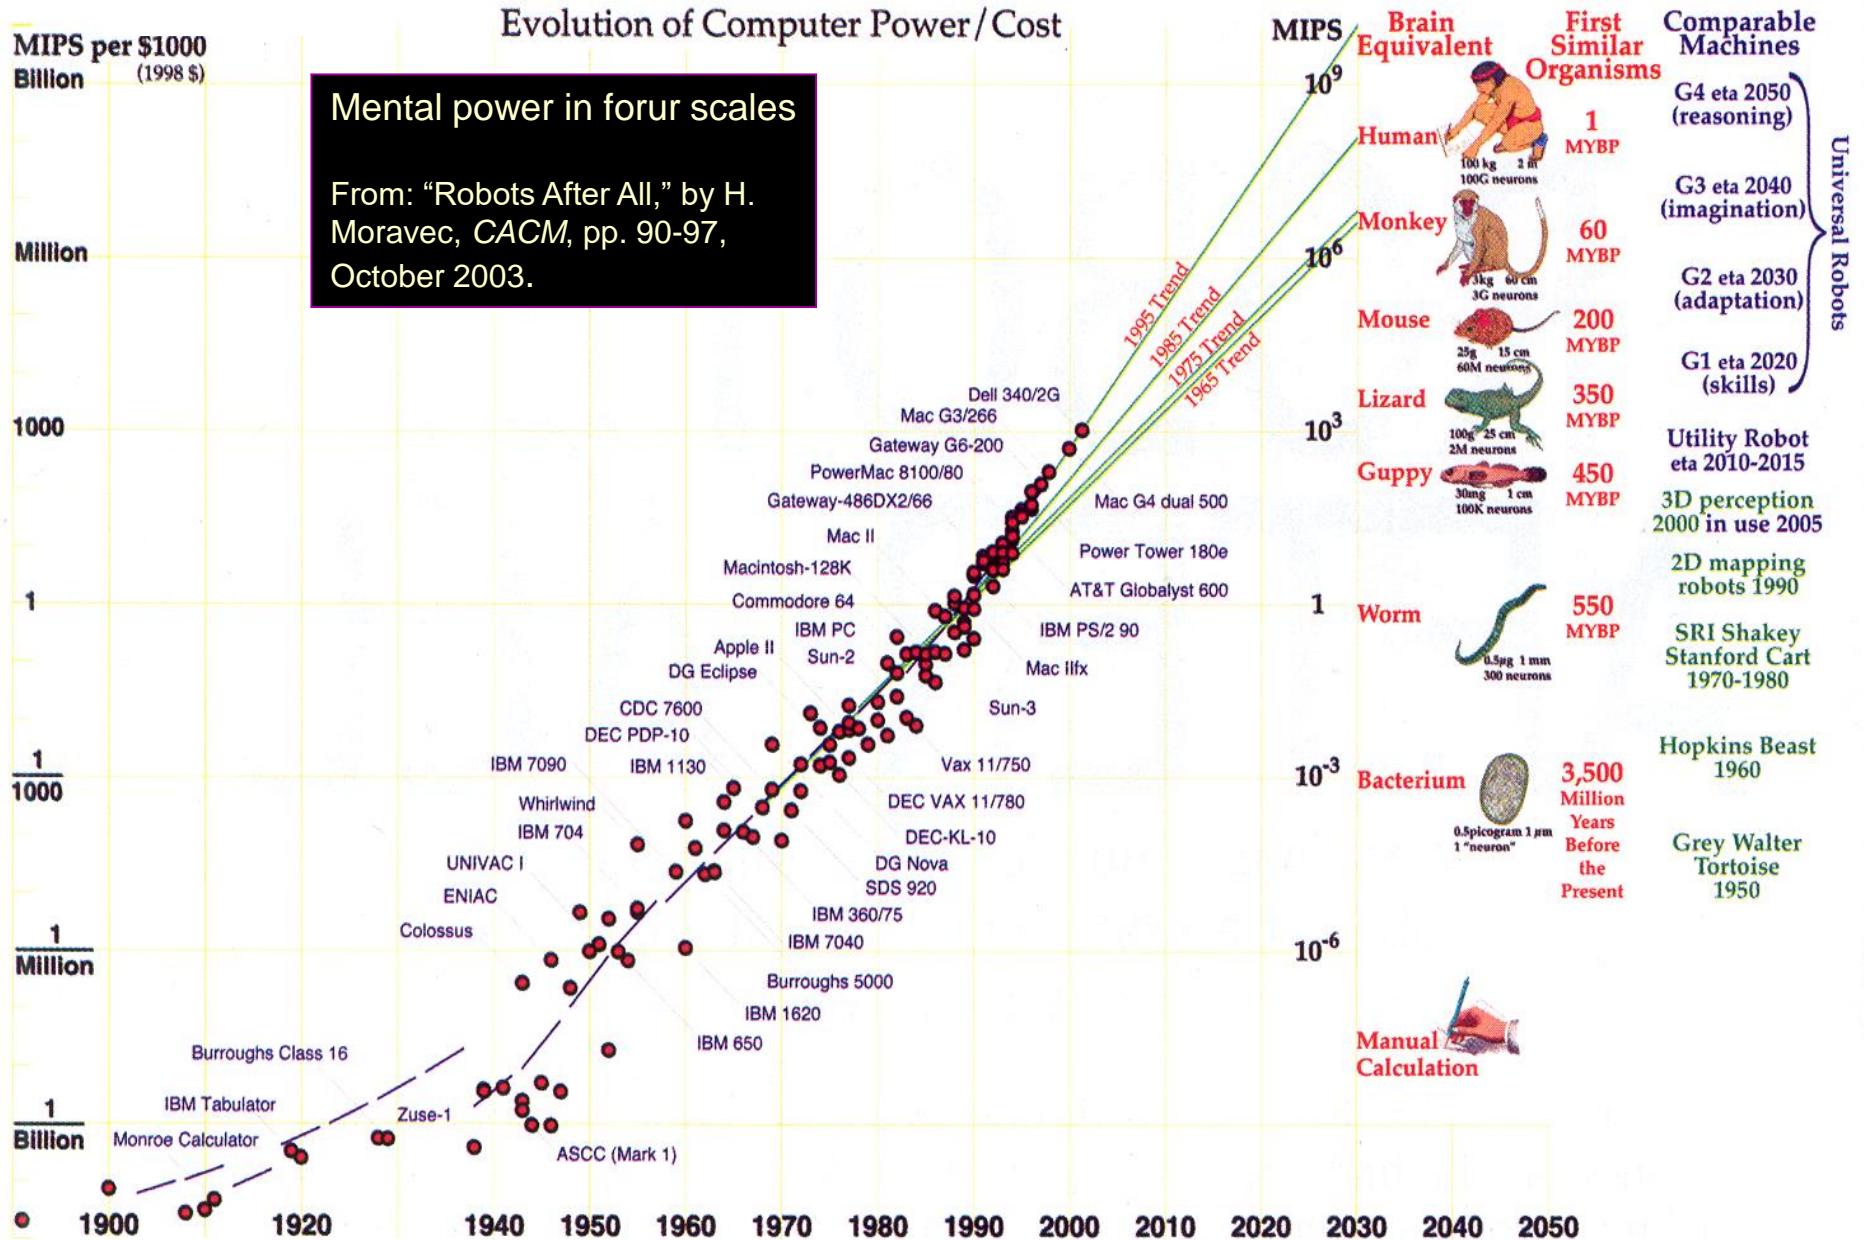
# *Why do we need big performance?*

- Higher speed (solve problems faster)
- Higher throughput (solve more problems)
- Higher computational power (solve larger problems)
- Examples:
  - Modelling
    - weather, environment, drugs, diseases
    - forecasting (earthquake, volcanoes, global warming)
    - Mechanics, chemistry
  - Simulations
    - design, production
    - bioinformatics, drug discovery
    - particle physics

# Evolution of Computer Power/Cost

## Mental power in four scales

From: "Robots After All," by H. Moravec, CACM, pp. 90-97, October 2003.





# *The Speed-of-Light Argument*



The speed of light is about 30 cm/ns.

Signals travel at 40-70% speed of light (say, 15 cm/ns).

If signals must travel 1.5 cm during the execution of an instruction, that instruction will take at least 0.1 ns; thus, performance will be limited to 10 GIPS.

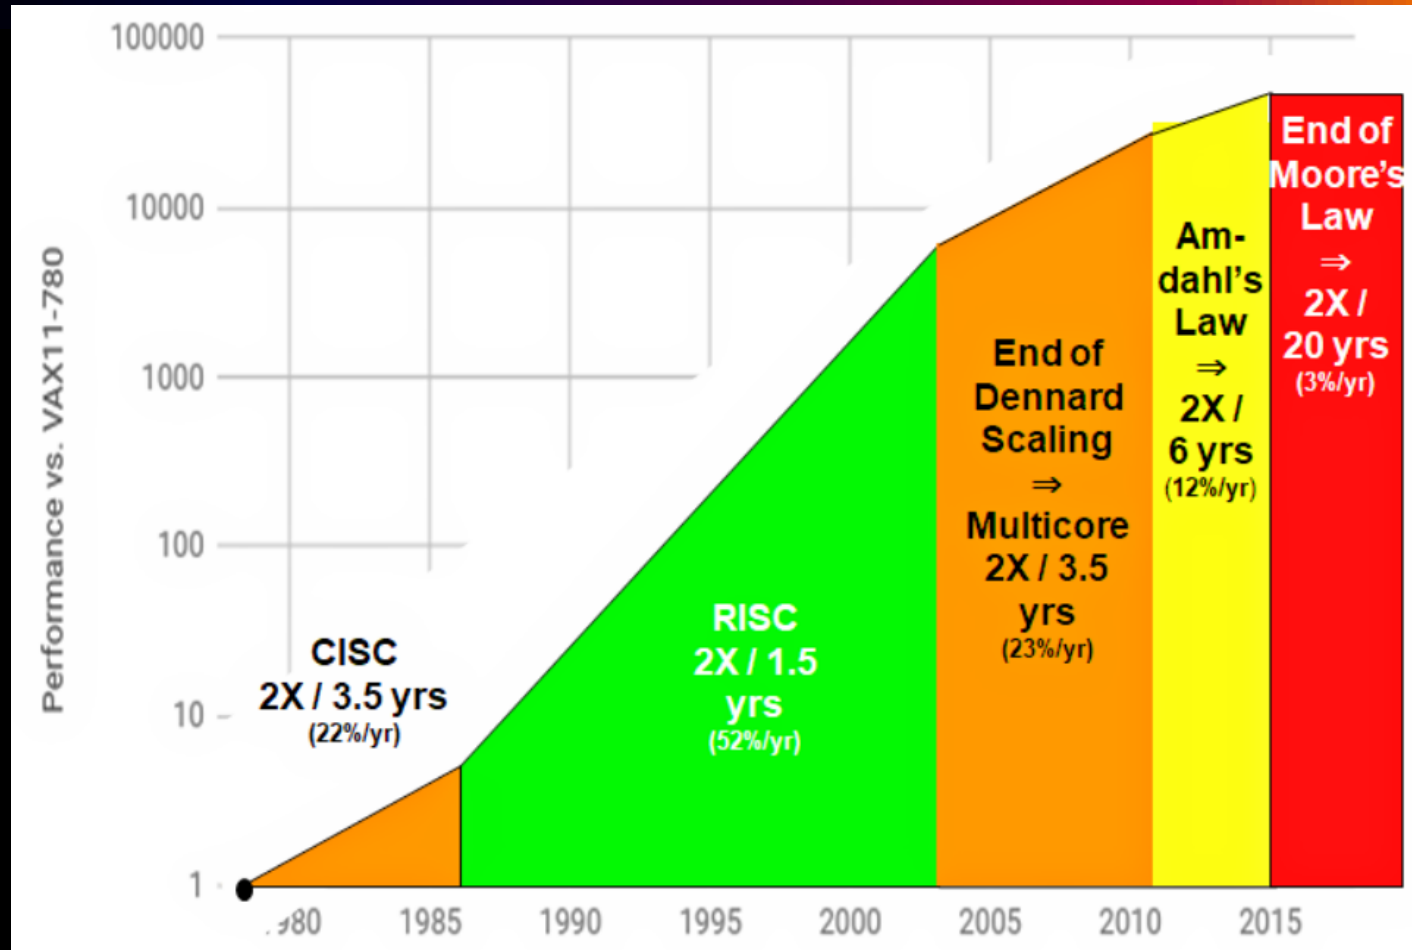
This limitation is eased by continued miniaturization, architectural methods such as cache memory, etc.; however, a fundamental limit does exist.

# *Performance development*



- Increase the clock frequency
  - Physical limits (3.7GHz, 9GHz+liquid Nitrogen cooling)
  - energy requirement (cubic dependency)
- Parallelization
  - overhead
  - hard to debug

# Processor Performance



Source: Hennessy and Patterson, *Computer Architecture: A Quantitative Approach*, 6/e. 2018

Reading: <https://www.eejournal.com/article/fifty-or-sixty-years-of-processor-developmentfor-this/>

# TOP 500 2013 november

| Rank | Site   | System   | Cores     | Rmax<br>(TFlop/s) | Rpeak<br>(TFlop/s) | Power<br>(kW) |
|------|--|--|-----------|-------------------|--------------------|---------------|
| 1    | National Super Computer Center in Guangzhou<br>China               | <b>Tianhe-2 (MilkyWay-2)</b> - TH-IVB-FEP Cluster, Intel Xeon E5-2692 12C 2.200GHz, TH Express-2, Intel Xeon Phi 31S1P<br>NUDT | 3,120,000 | 33,862.7          | 54,902.4           | 17,808        |
| 2    | DOE/SC/Oak Ridge National Laboratory<br>United States              | <b>Titan</b> - Cray XK7 , Opteron 6274 16C 2.200GHz, Cray Gemini interconnect, NVIDIA K20x<br>Cray Inc.                        | 560,640   | 17,590.0          | 27,112.5           | 8,209         |
| 3    | DOE/NNSA/LLNL<br>United States                                     | <b>Sequoia</b> - BlueGene/Q, Power BQC 16C 1.60 GHz, Custom<br>IBM   | 1,572,864 | 17,173.2          | 20,132.7           | 7,890         |
| 4    | RIKEN Advanced Institute for Computational Science (AICS)<br>Japan | K computer, SPARC64 VIIIfx 2.0GHz, Tofu interconnect<br>Fujitsu  | 705,024   | 10,510.0          | 11,280.4           | 12,660        |
| 5    | DOE/SC/Argonne National Laboratory<br>United States                | <b>Mira</b> - BlueGene/Q, Power BQC 16C 1.60GHz, Custom<br>IBM   | 786,432   | 8,586.6           | 10,066.3           | 3,945         |
| 6    | Swiss National Supercomputing Centre (CSCS)<br>Switzerland         | <b>Piz Daint</b> - Cray XC30, Xeon E5-2670 8C 2.600GHz, Aries interconnect , NVIDIA K20x<br>Cray Inc.                          | 115,984   | 6,271.0           | 7,788.9            | 2,325         |
| 7    | Texas Advanced Computing Center/Univ. of Texas<br>United States    | <b>Stampede</b> - PowerEdge C8220, Xeon E5-2680 8C 2.700GHz, Infiniband FDR, Intel Xeon Phi SE10P<br>Dell                      | 462,462   | 5,168.1           | 8,520.1            | 4,510         |
| 8    | Forschungszentrum Juelich (FZJ)<br>Germany                         | <b>JUQUEEN</b> - BlueGene/Q, Power BQC 16C 1.600GHz, Custom Interconnect<br>IBM  | 458,752   | 5,008.9           | 5,872.0            | 2,301         |
| 9    | DOE/NNSA/LLNL<br>United States                                     | <b>Vulcan</b> - BlueGene/Q, Power BQC 16C 1.600GHz, Custom Interconnect<br>IBM   | 393,216   | 4,293.3           | 5,033.2            | 1,972         |
| 10   | Leibniz Rechenzentrum  | <b>SuperMUC</b> - iDataPlex DX360M4, Xeon E5-2680 8C   | 147,456   | 2,897.0           | 3,185.1            | 3,423         |

Only the  
6th is  
European

# TOP 500 2017 june

| Rank | System   | Cores      | Rmax<br>(TFlop/s) | Rpeak<br>(TFlop/s) | Power<br>(kW) |
|------|--|------------|-------------------|--------------------|---------------|
| 1    | <b>Sunway TaihuLight</b> - Sunway MPP, Sunway SW26010 260C<br>1.45GHz, Sunway , NRCPC<br>National Supercomputing Center in Wuxi<br>China   | 10,649,600 | 93,014.6          | 125,435.9          | 15,371        |
| 2    | <b>Tianhe-2 (MilkyWay-2)</b> - TH-IVB-FEP Cluster, Intel Xeon E5-2692<br>12C 2.200GHz, TH Express-2, Intel Xeon Phi 31S1P , NUDT<br>National Super Computer Center in Guangzhou<br>China | 3,120,000  | 33,862.7          | 54,902.4           | 17,808        |
| 3    | <b>Piz Daint</b> - Cray XC50, Xeon E5-2690v3 12C 2.6GHz, Aries<br>interconnect , NVIDIA Tesla P100 , Cray Inc.<br>Swiss National Supercomputing Centre (CSCS)<br>Switzerland             | 361,760    | 19,590.0          | 25,326.3           | 2,272         |
| 4    | <b>Titan</b> - Cray XK7, Opteron 6274 16C 2.200GHz, Cray Gemini<br>interconnect, NVIDIA K20x , Cray Inc.<br>DOE/SC/Oak Ridge National Laboratory<br>United States                        | 560,640    | 17,590.0          | 27,112.5           | 8,209         |
| 5    | <b>Sequoia</b> - BlueGene/Q, Power BQC 16C 1.60 GHz, Custom , IBM<br>DOE/NNSA/LLNL<br>United States  | 1,572,864  | 17,173.2          | 20,132.7           | 7,890         |
| 6    | <b>Cori</b> - Cray XC40, Intel Xeon Phi 7250 68C 1.4GHz, Aries<br>interconnect , Cray Inc.<br>DOE/SC/LBNL/NERSC<br>United States   | 622,336    | 14,014.7          | 27,880.7           | 3,939         |

Now the  
3rd is  
European

# TOP 500 2018 june

| Rank | System   | Cores      | Rmax<br>(TFlop/s) | Rpeak<br>(TFlop/s) | Power<br>(kW) |
|------|--|------------|-------------------|--------------------|---------------|
| 1    | <b>Summit</b> - IBM Power System AC922, IBM POWER9 22C 3.07GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband , IBM DOE/SC/Oak Ridge National Laboratory United States   | 2,282,544  | 122,300.0         | 187,659.3          | 8,806         |
| 2    | <b>Sunway TaihuLight</b> - Sunway MPP, Sunway SW26010 260C 1.45GHz, Sunway , NRCPC National Supercomputing Center in Wuxi China  | 10,649,600 | 93,014.6          | 125,435.9          | 15,371        |
| 3    | <b>Sierra</b> - IBM Power System S922LC, IBM POWER9 22C 3.1GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband , IBM DOE/NNSA/LLNL United States  | 1,572,480  | 71,610.0          | 119,193.6          |               |
| 4    | <b>Tianhe-2A</b> - TH-IVB-FEP Cluster, Intel Xeon E5-2692v2 12C 2.2GHz, TH Express-2, Matrix-2000 , NUDT National Super Computer Center in Guangzhou China   | 4,981,760  | 61,444.5          | 100,678.7          | 18,482        |
| 5    | <b>AI Bridging Cloud Infrastructure (ABCI)</b> - PRIMERGY CX2550 M4, Xeon Gold 6148 20C 2.4GHz, NVIDIA Tesla V100 SXM2, Infiniband EDR , Fujitsu National Institute of Advanced Industrial Science and Technology (AIST) Japan | 391,680    | 19,880.0          | 32,576.6           | 1,649         |
| 6    | <b>Piz Daint</b> - Cray XC50, Xeon E5-2690v3 12C 2.6GHz, Aries interconnect , NVIDIA Tesla P100 , Cray Inc. Swiss National Supercomputing Centre (CSCS) Switzerland  | 361,760    | 19,590.0          | 25,326.3           | 2,272         |

The 6th  
again.

# TOP 500 2020 june

| Rank | System  | Cores      | Rmax<br>(TFlop/s) | Rpeak<br>(TFlop/s) | Power<br>(kW) |
|------|---|------------|-------------------|--------------------|---------------|
| 1    | <b>Supercomputer Fugaku</b> - Supercomputer Fugaku,<br>A64FX 48C 2.2GHz, Tofu interconnect D, Fujitsu<br>RIKEN Center for Computational Science<br>Japan                                  | 7,299,072  | 415,530.0         | 513,854.7          | 28,335        |
| 2    | <b>Summit</b> - IBM Power System AC922, IBM POWER9 22C<br>3.07GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR<br>Infiniband, IBM<br>DOE/SC/Oak Ridge National Laboratory<br>United States | 2,414,592  | 148,600.0         | 200,794.9          | 10,096        |
| 3    | <b>Sierra</b> - IBM Power System AC922, IBM POWER9 22C<br>3.1GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR<br>Infiniband, IBM / NVIDIA / Mellanox<br>DOE/NNSA/LLNL<br>United States     | 1,572,480  | 94,640.0          | 125,712.0          | 7,438         |
| 4    | <b>Sunway TaihuLight</b> - Sunway MPP, Sunway SW26010<br>260C 1.45GHz, Sunway, NRCP<br>National Supercomputing Center in Wuxi<br>China  | 10,649,600 | 93,014.6          | 125,435.9          | 15,371        |
| 5    | <b>Tianhe-2A</b> - TH-IVB-FEP Cluster, Intel Xeon E5-2692v2<br>12C 2.2GHz, TH Express-2, Matrix-2000, NUDT<br>National Super Computer Center in Guangzhou<br>China                        | 4,981,760  | 61,444.5          | 100,678.7          | 18,482        |
| 6    | <b>HPC5</b> - PowerEdge C4140, Xeon Gold 6252 24C 2.1GHz,<br>NVIDIA Tesla V100, Mellanox HDR Infiniband, Dell EMC<br>Eni S.p.A.<br>Italy  | 669,760    | 35,450.0          | 51,720.8           | 2,252         |

The 6th  
again.

# TOP 500 2021 june

| Rank | System  | Cores      | Rmax<br>(TFlop/s) | Rpeak<br>(TFlop/s) | Power<br>(kW) |
|------|---|------------|-------------------|--------------------|---------------|
| 1    | <b>Supercomputer Fugaku</b> - Supercomputer Fugaku, A64FX 48C 2.2GHz, Tofu interconnect D, Fujitsu<br>RIKEN Center for Computational Science<br>Japan                               | 7,630,848  | 442,010.0         | 537,212.0          | 29,899        |
| 2    | <b>Summit</b> - IBM Power System AC922, IBM POWER9 22C 3.07GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband, IBM<br>DOE/SC/Oak Ridge National Laboratory<br>United States | 2,414,592  | 148,600.0         | 200,794.9          | 10,096        |
| 3    | <b>Sierra</b> - IBM Power System AC922, IBM POWER9 22C 3.1GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband, IBM / NVIDIA / Mellanox<br>DOE/NNSA/LLNL<br>United States     | 1,572,480  | 94,640.0          | 125,712.0          | 7,438         |
| 4    | <b>Sunway TaihuLight</b> - Sunway MPP, Sunway SW26010 260C 1.45GHz, Sunway, NRCPC<br>National Supercomputing Center in Wuxi<br>China  | 10,649,600 | 93,014.6          | 125,435.9          | 15,371        |
| 5    | <b>Perlmutter</b> - HPE Cray EX235n, AMD EPYC 7763 64C 2.45GHz, NVIDIA A100 SXM4 40 GB, Slingshot-10, HPE<br>DOE/SC/LBNL/NERSC<br>United States                                     | 706,304    | 64,590.0          | 89,794.5           | 2,528         |
| 6    | <b>Selene</b> - NVIDIA DGX A100, AMD EPYC 7742 64C 2.25GHz, NVIDIA A100, Mellanox HDR Infiniband, Nvidia<br>NVIDIA Corporation<br>United States                                     | 555,520    | 63,460.0          | 79,215.0           | 2,646         |

Now  
only the  
8th is  
European



# TOP 500 2022 june

| Rank | System  | Cores      | Rmax<br>(PFlop/s) | Rpeak<br>(PFlop/s) | Power<br>(kW) |
|------|---|------------|-------------------|--------------------|---------------|
| 1    | <b>Frontier</b> - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE<br>DOE/SC/Oak Ridge National Laboratory<br>United States      | 8,730,112  | 1,102.00          | 1,685.65           | 21,100        |
| 2    | <b>Supercomputer Fugaku</b> - Supercomputer Fugaku, A64FX 48C 2.2GHz, Tofu interconnect D, Fujitsu<br>RIKEN Center for Computational Science<br>Japan                               | 7,630,848  | 442.01            | 537.21             | 29,899        |
| 3    | <b>LUMI</b> - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE<br>EuroHPC/CSC<br>Finland   | 1,110,144  | 151.90            | 214.35             | 2,942         |
| 4    | <b>Summit</b> - IBM Power System AC922, IBM POWER9 22C 3.07GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband, IBM<br>DOE/SC/Oak Ridge National Laboratory<br>United States | 2,414,592  | 148.60            | 200.79             | 10,096        |
| 5    | <b>Sierra</b> - IBM Power System AC922, IBM POWER9 22C 3.1GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR Infiniband, IBM / NVIDIA / Mellanox<br>DOE/NNSA/LLNL<br>United States     | 1,572,480  | 94.64             | 125.71             | 7,438         |
| 6    | <b>Sunway TaihuLight</b> - Sunway MPP, Sunway SW26010 260C 1.45GHz, Sunway, NRCPC<br>National Supercomputing Center in Wuxi<br>China  | 10,649,600 | 93.01             | 125.44             | 15,371        |

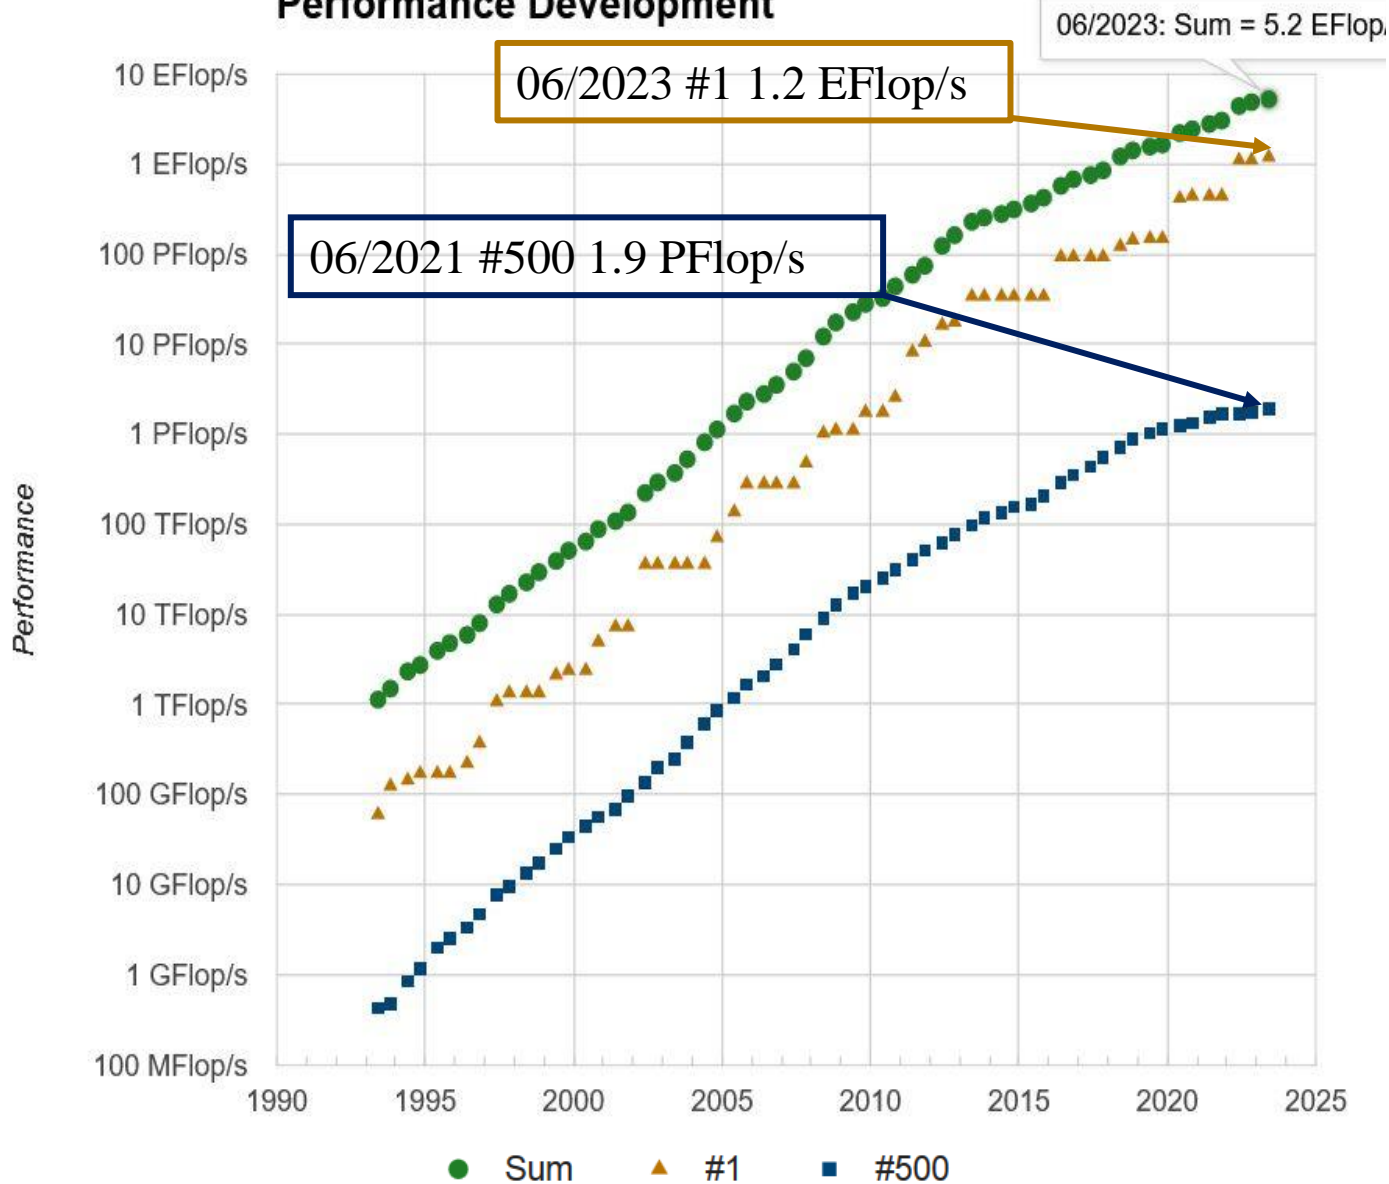
Now the  
3th is  
European

# TOP 500 2023 june

| Rank | System   | Cores     | Rmax<br>(PFlop/s) | Rpeak<br>(PFlop/s) | Power<br>(kW) |
|------|--|-----------|-------------------|--------------------|---------------|
| 1    | <b>Frontier</b> - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE<br>DOE/SC/Oak Ridge National Laboratory<br>United States         | 8,699,904 | 1,194.00          | 1,679.82           | 22,703        |
| 2    | <b>Supercomputer Fugaku</b> - Supercomputer Fugaku, A64FX 48C 2.2GHz, Tofu interconnect D, Fujitsu<br>RIKEN Center for Computational Science<br>Japan                                  | 7,630,848 | 442.01            | 537.21             | 29,899        |
| 3    | <b>LUMI</b> - HPE Cray EX235a, AMD Optimized 3rd Generation EPYC 64C 2GHz, AMD Instinct MI250X, Slingshot-11, HPE<br>EuroHPC/CSC<br>Finland  | 2,220,288 | 309.10            | 428.70             | 6,016         |
| 4    | <b>Leonardo</b> - BullSequana XH2000, Xeon Platinum 8358 32C 2.6GHz, NVIDIA A100 SXM4 64 GB, Quad-rail<br>NVIDIA HDR100 Infiniband, Atos<br>EuroHPC/CINECA<br>Italy                    | 1,824,768 | 238.70            | 304.47             | 7,404         |
| 5    | <b>Summit</b> - IBM Power System AC922, IBM POWER9 22C 3.07GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR<br>Infiniband, IBM<br>DOE/SC/Oak Ridge National Laboratory<br>United States | 2,414,592 | 148.60            | 200.79             | 10,096        |
| 6    | <b>Sierra</b> - IBM Power System AC922, IBM POWER9 22C 3.1GHz, NVIDIA Volta GV100, Dual-rail Mellanox EDR<br>Infiniband, IBM / NVIDIA / Mellanox<br>DOE/NNSA/LLNL<br>United States     | 1,572,480 | 94.64             | 125.71             | 7,438         |

Now the  
3th and  
the 4th is  
European

## Performance Development



# *Behind the development*

- The density of element
  - 1971: 10 $\mu$ m, 1984: 1 $\mu$ m, 2001: 130nm, 2020: 5nm, 2022: 3nm, 2024: 2nm
  - [https://en.wikipedia.org/wiki/2\\_nm\\_process](https://en.wikipedia.org/wiki/2_nm_process)
  - <https://youtu.be/YIkMaQJSyP8>
- Development of network technology
- Development of storage technology
- Development of memory
- Many cores technology

# *Overview of parallel programming*

## **Why is parallel processing is important?**

- For greater performance? Only for performance?
  - Parallelism in this case is just technology
  - not important for the application designer
  - Must be covered (such as hardware registers, cache, ...)
- A possible tool for modelling reality
- Simplify design work

## **Problems with parallel processing**

- It is often harder to manage than serial processing
- fragile (non deterministic behaviour)
- deadlock problems can occur
- resource allocation issues can occur
- Error detection is not easy

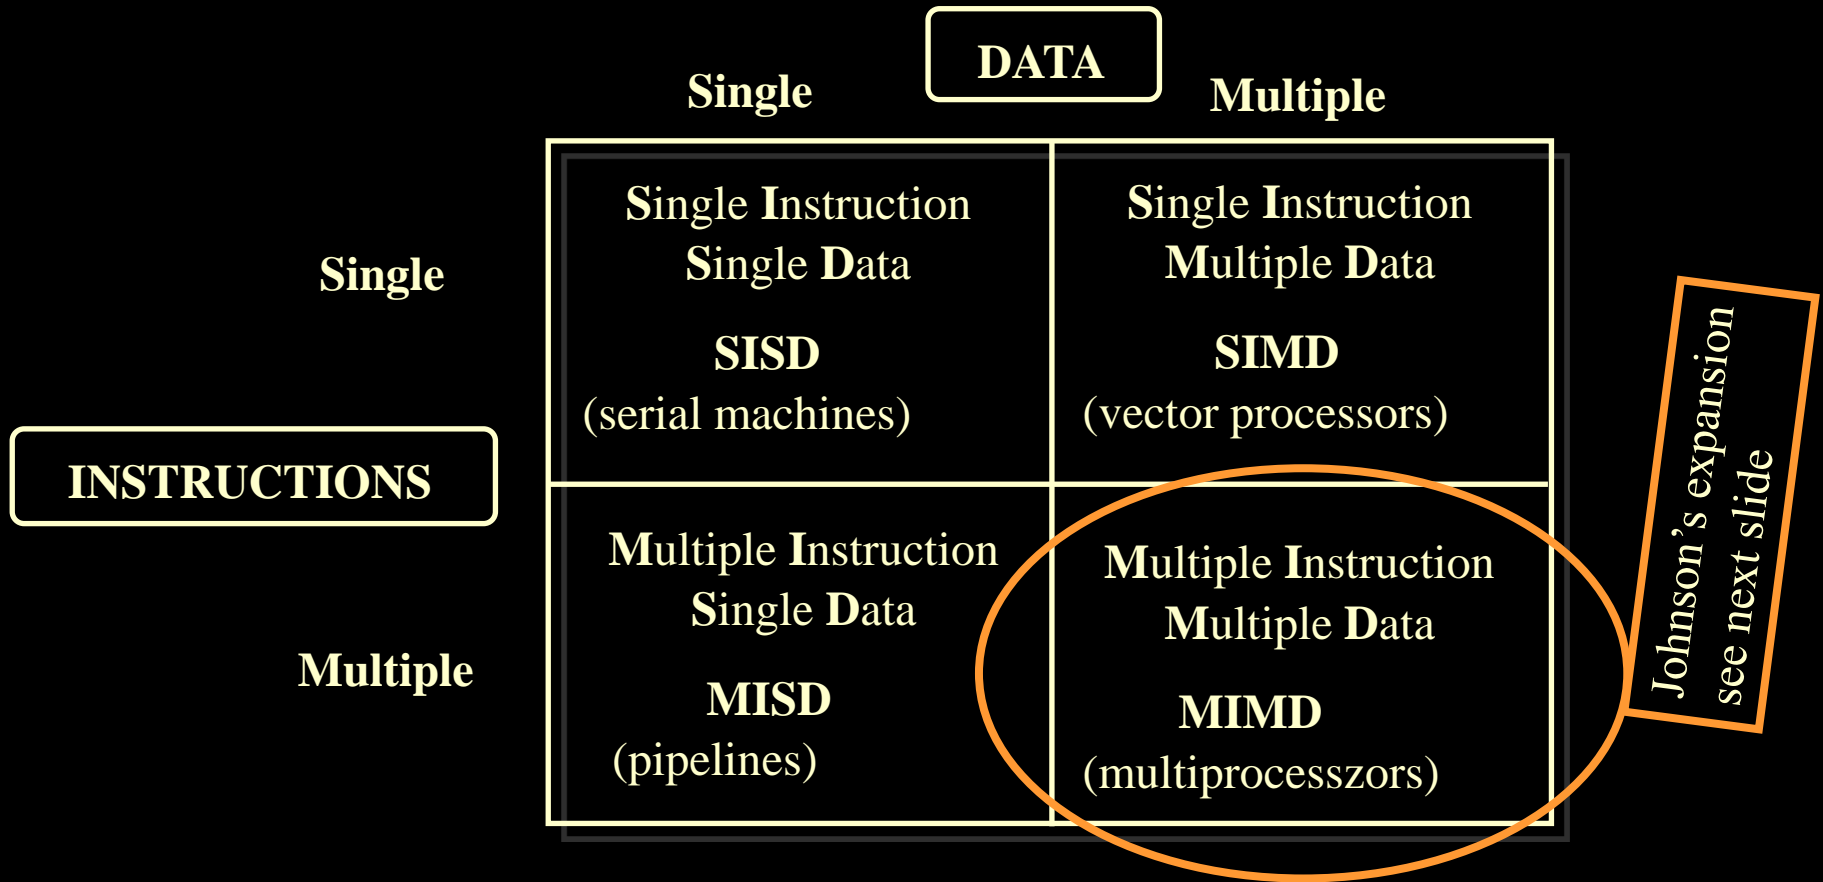
**These are real experiences, but not a necessity!**

# *Parallel machine models*

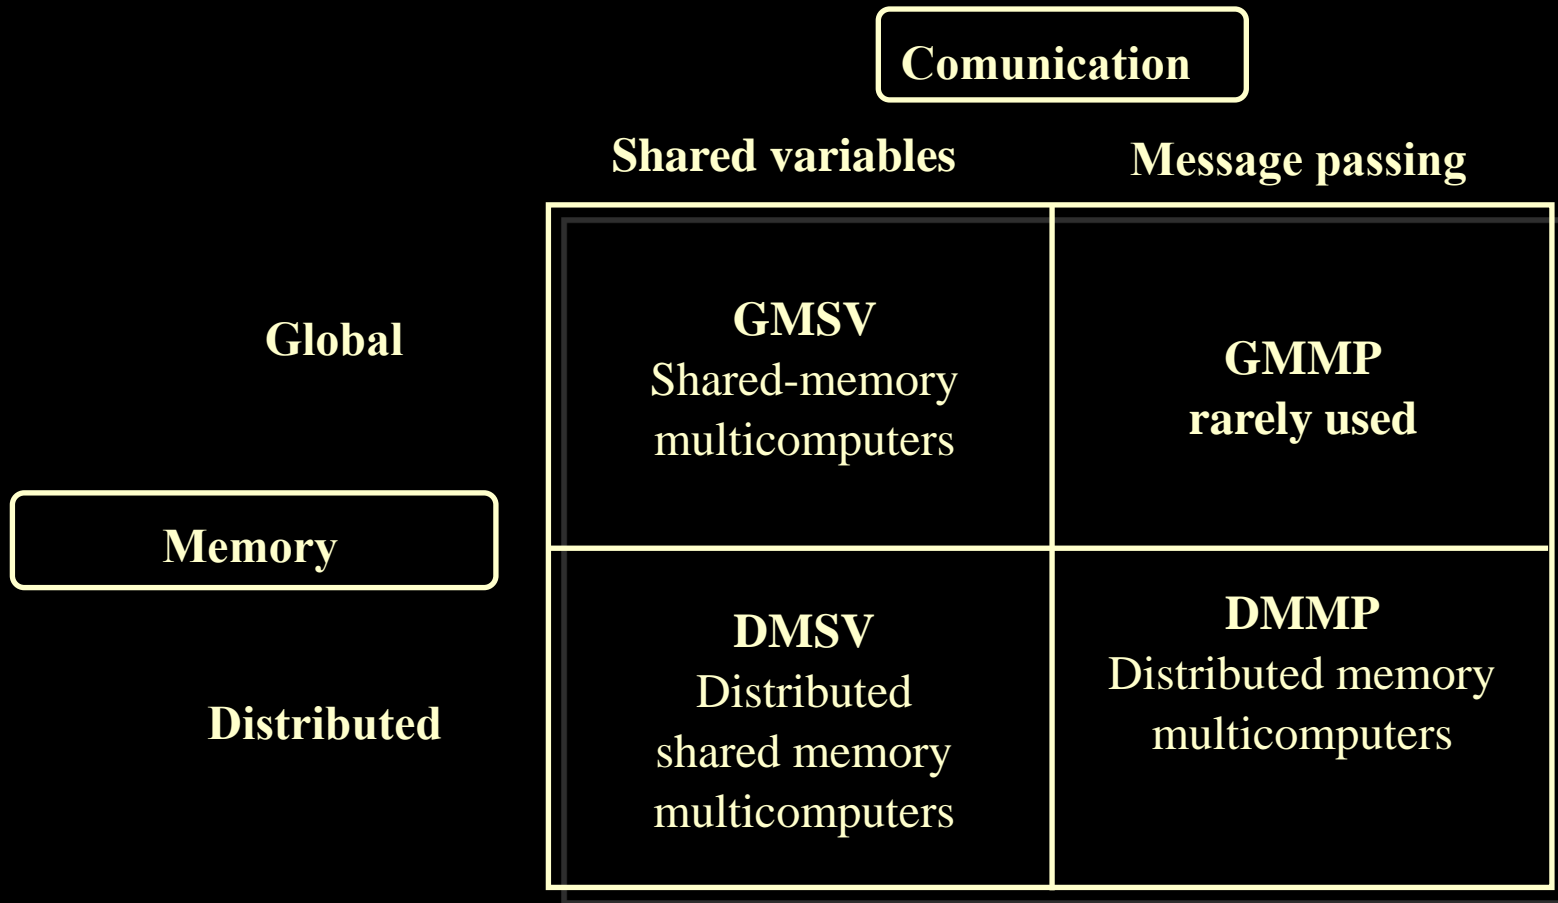
---

- Several models have emerged.
- The simplest is Flynn's model (taxonomy), which considers the parallel machine as the extension of the Neumann model.
- Another commonly used model is the idealized parallel computer model

# *Flynn's taxonomy*

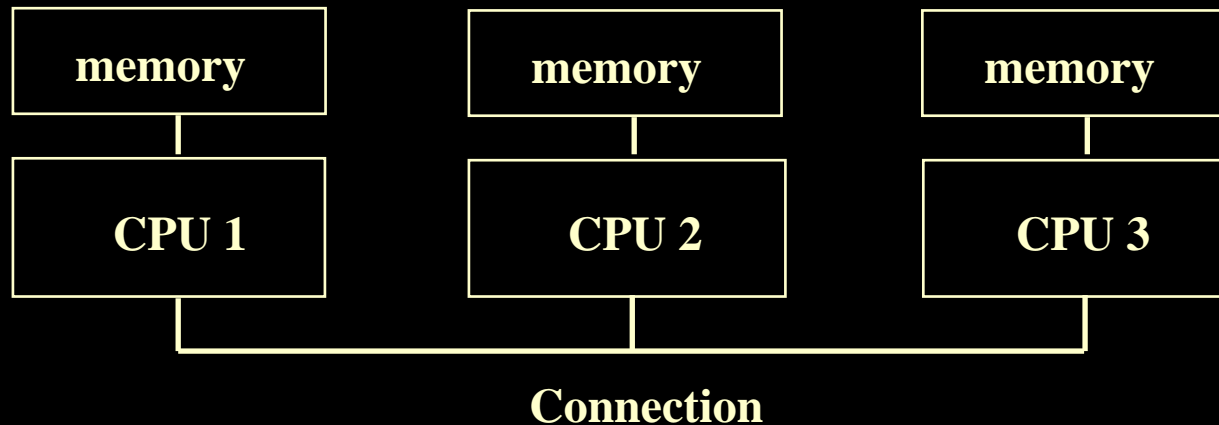


# *Johnson's expansion of MIMD*





# *Idealized parallel computer*

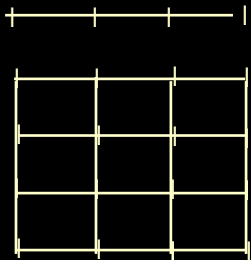


- Several processors are working on the same problem.
- Each processor has its own memory and address space.
- They coordinate with messages and can handover data.
- Local memory is faster to access.
- Transmission speed is independent of channel traffic.

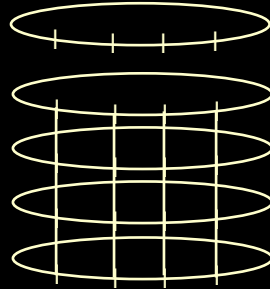
# *Features of architectures*

- Distribution of processors
- Homogeneous or heterogeneous
- Interconnect bandwidth and latency
- Topology

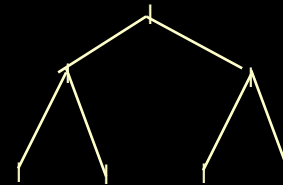
Nets



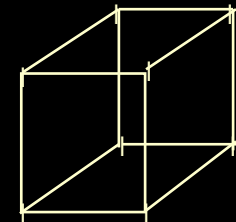
Rings



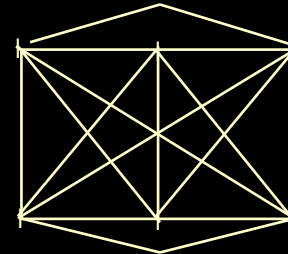
Trees



Hibercubs



Fully connected



# *Programming model*

---

- Shared memory
- Distributed shared memory
- Message passing

In fact, none of the models is closely tied to the actual physical architecture

# *Features of shared memory principle*



- Benefits of using shared memory
  - Unified access
  - Simplified programming
  - Depending on the characteristics of hw, good speed-up ratings are available
- Disadvantages of using shared memory
  - Memory access may be a bottleneck
  - Not well scalable
  - Solving cache problems requires separate hardware
  - Debugging is difficult

# *Features of the distr. mem. principle*

- **Benefits of using distributed memory**
  - Scalable
  - cost-effective
  - Increasing redundancy can increase reliability
  - It also works with special tools, accelerators
- **Disadvantages of using distributed memory**
  - Demands a higher communication
  - Not all algorithms can be parallelised this way
  - Existing serial programs and applications using shared memory must be redrawn
  - Good speed ups are difficult to achieve
  - Debugging problems

# *Classes of parallel machines*

- Machines with Vector processors, array proc.
  - optimized operations for vector data
  - one operating system instance
- Symmetric Multiprocessor (SMP)
  - many of same processor have common memory
  - one operating system instance
  - NUMA, ccNUMA
- Massively Parallel (MPP)
  - many processors with a fast internal network
  - distributed memory
  - there are many instances of operating system
- Cluster
  - many machines connected with a fast (commerce) network
  - distributed memory
  - in many instances it may be a heterogeneous operating system

# Performance Measurement

## Speed Up:

$$S_n = T_s / T_n$$

where:  $S_n$  speed up gained with N processor  
 $T_s$  run time in Serial Execution  
 $T_n$  run time in case of N processor

## Efficiency:

$$E_n = S_n / N$$

where:  $E_n$  efficiency gained with N processor  
 $S_n$  speed up gained with N processor  
N number of processors

## Redundancy (redundancy):

$$r = C_p / C_s$$

where: r is the redundancy of the parallel program  
 $C_p$  The number of operations in the parallel program  
 $C_s$  The number of operations in the serial program

# Speed Up limit

- **Amdahl's upper limit:**  $S_a = 1 / (s + (1-s) / N)$

where:  $S_a$  the upper limit of the speed up which can be gained by  $N$  processors

$s$  part of the task that can NOT be parallelized

$N$  Number of processors

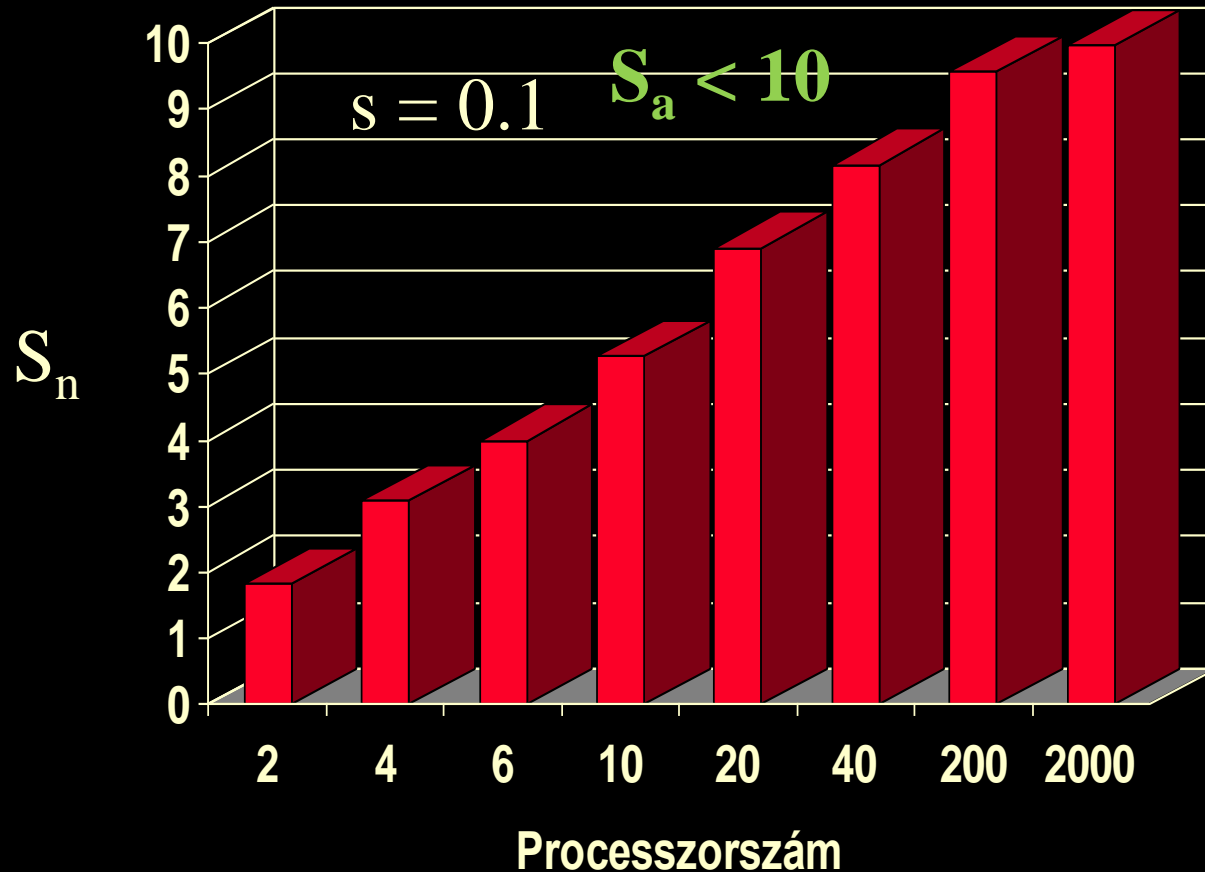
**Leaving  $(1-s) / N$ :**  $S_a < 1 / s$

gives the above context which provides an upper limit.

For example, 10% non-parallelizable part gives  $1 / 0.1 = 10$  for the upper limit of speed up.



# *It can not be accelerated unlimited*



# *Programming languages*

- Linda - a shared memory model, Tuple Space
  - 1986 – (C, later: Java, Python, Ruby, Lua ...)
  - out - copies a tuple to the shared area
  - in, in<sub>p</sub> - inports a tuple (will be removed)
  - rd, rd<sub>p</sub> – reads a tuple (<sub>p</sub> – probe, non blocking)
  - eval - executes a function as a process

In fact, there is a language supplement for a simple model but there are implementation difficulties, especially in message passing architectures.

# *Linda example #1 (hello)*

```
#define NPROC 8
int hello(int id) {
    int j;
    in("count", ?j);
    printf("Hello World from proc%d: %d\n", id, j);
    return j+1;
}
int real_main() {
    int i;
    out("count", 0);
    for (i = 0; i < NPROC; ++i)
        eval("count", hello(i));
    in("count", NPROC);
    printf("All processors done\n");
    return 0;
}
```

## *Linda example #2 (queue)*

```
int init() {  
    out("head", 0);  
    out("tail", 0);  
}
```

```
void push(int elem) {  
    int tail;  
    in ("tail", ?tail);  
    out("elem", tail, elem);  
    out("tail", tail+1);  
}  
  
int pop() {  
    int head, elem;  
    in ("head", ?head);  
    in ("elem", head, ?elem);  
    out("head", head+1);  
    return elem;  
}
```

# *Programming languages/2*

- Express - distributed memory model using message passing technique. Portable, supports logical connection of nodes.

~160 Routine callable from C and Fortran:

- start, shutdown
- structure of logical communication topology
- communication between programs,
- Synchronization
- file operations
- graphic operations
- performance analysis, debugging

# *Programming languages/3*

---

- PVM – (Parallel Virtual Machine)  
Distributed Memory Model
  - 1989 – 1993
- ~70 Routines callable from C and Fortran:
  - start, shutdown
  - communication between programs, synchronization

# *What is PVM used for?*



- The Poor Man's supercomputer

free CPU capacities can be collected from  
workstations and PCs

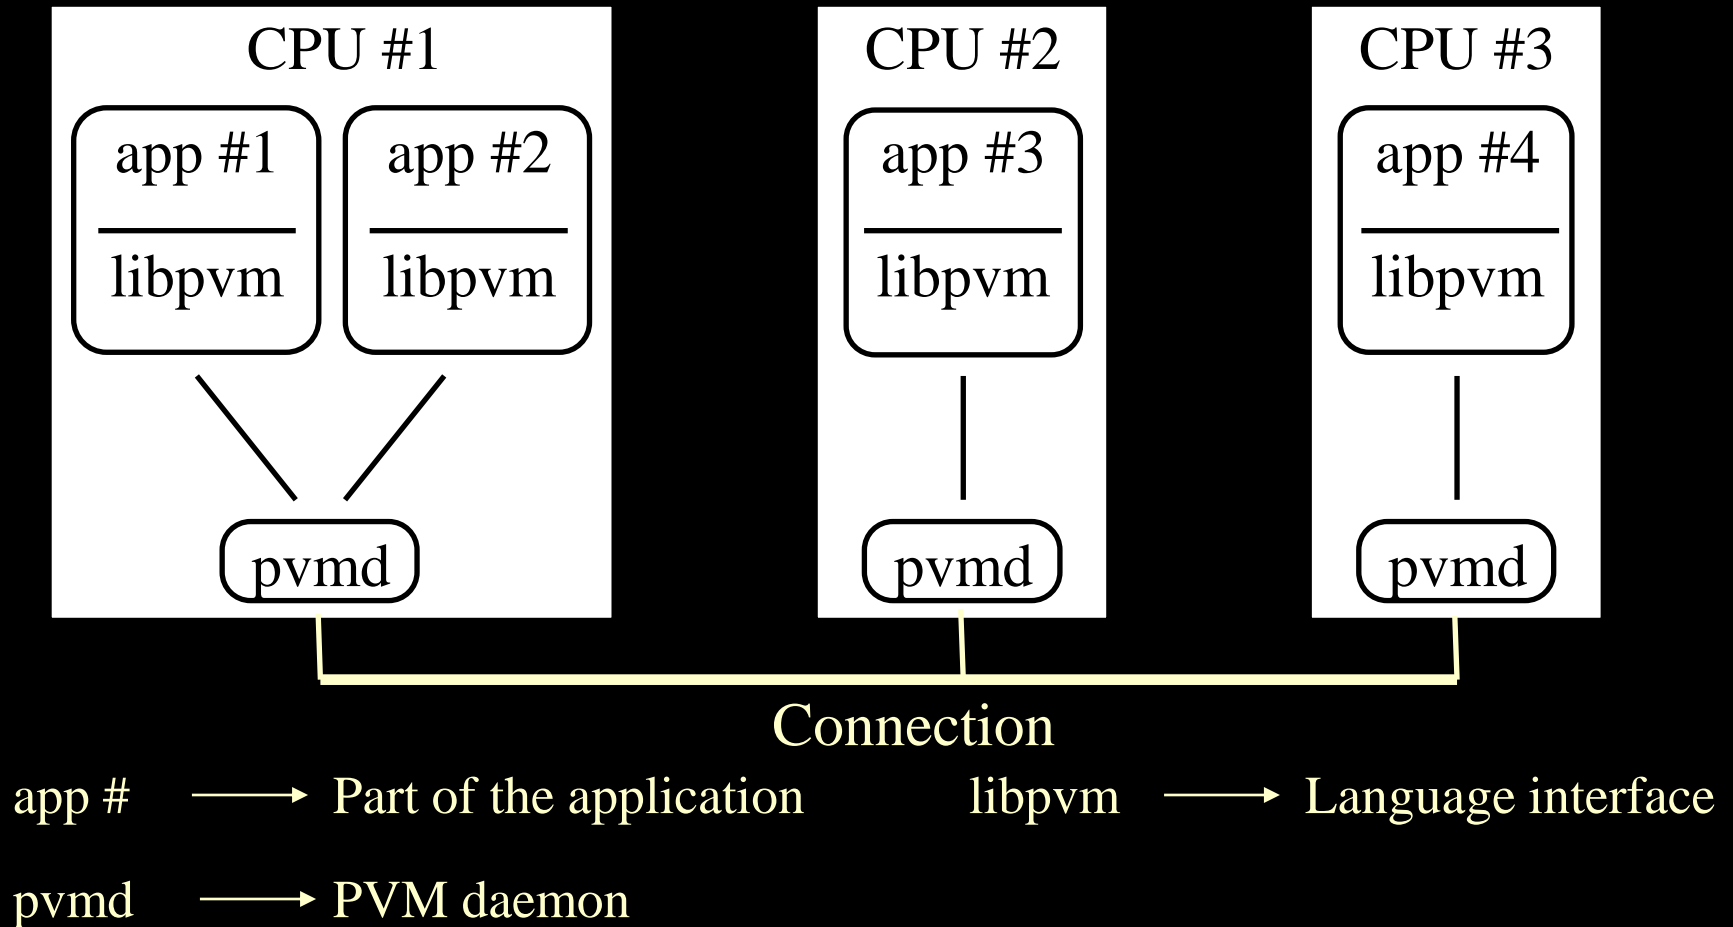
- Combining multiple supercomputers can create incredible  
computing capacity

- Educational tool

It is an effective tool for teaching parallel  
programming

- Research tool  
scalable and cost-effective

# *The basic concept of PVM*





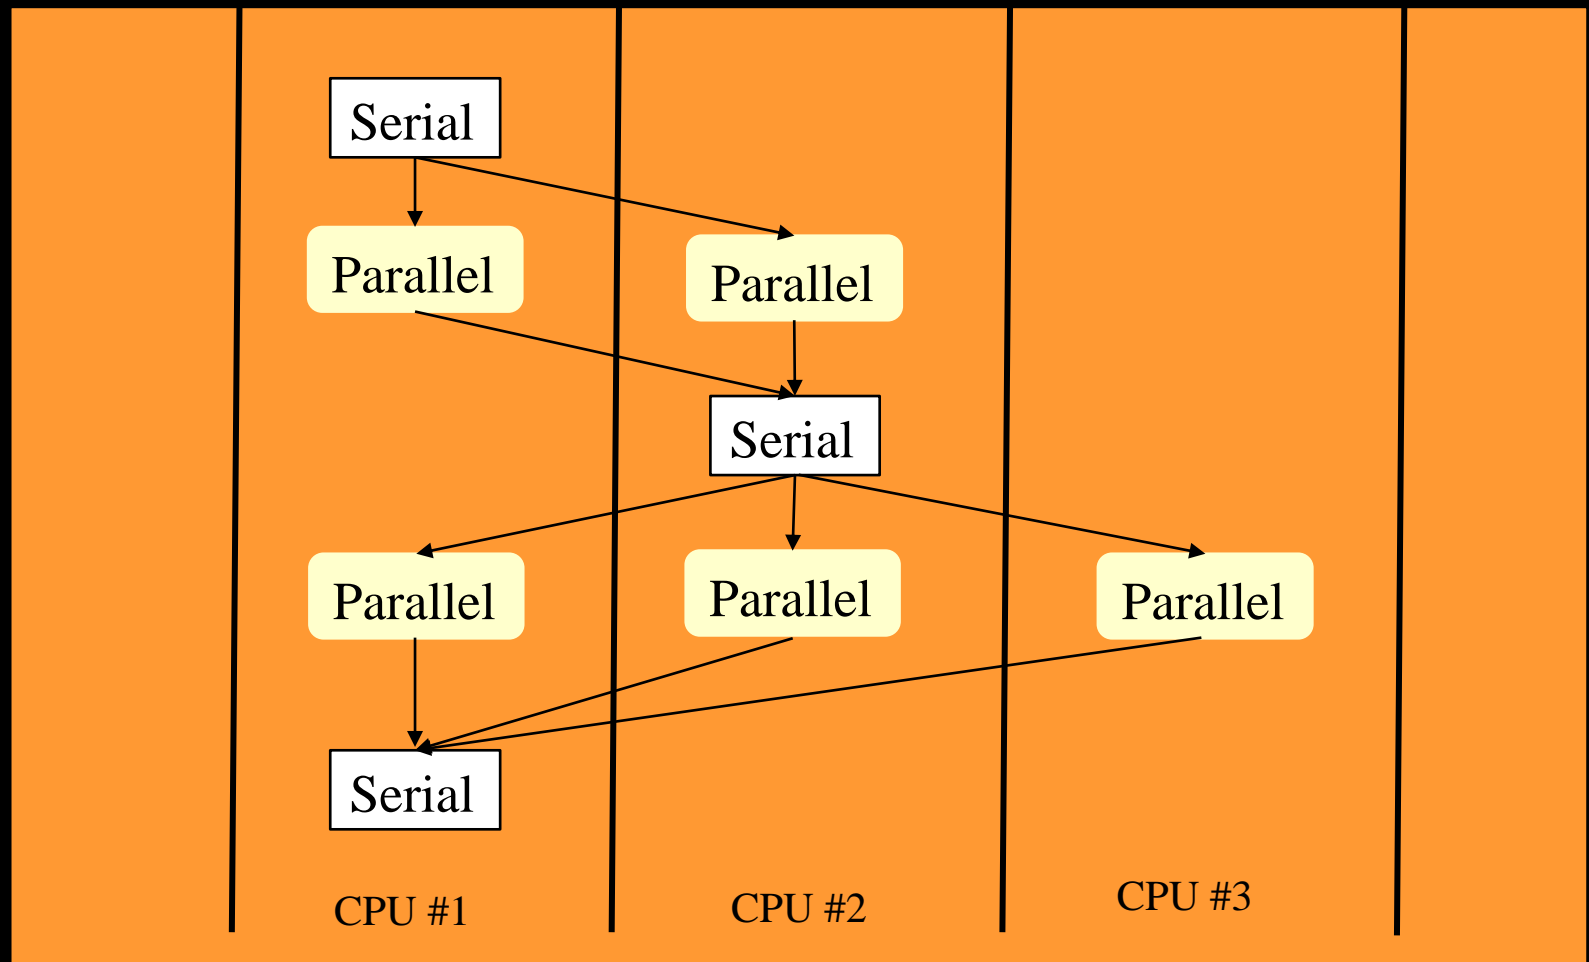
# *libpvm*



**The features provided by libpvm can be divided into four groups:**

- Administrative functions  
virtual machine start, stop, add a new node, status query,
- Process Manager Features  
Launching and stopping processes
- Data Transmission Features  
Composing, sending, receiving, packing messages
- Synchronization features  
By sending a message or using a barrier

# *Structure of a parallel program*



# Parallelization strategies

## Embarrassingly parallel

- We run the serial version of the program in parallel with different data.
- It is only a satisfactory method if the serial version has a tolerable running time.

## Parallelization of cycles

- It can be applied if each iteration is independent of each other

## Divide & conquer (master / slave)

- A supervisor task runs on one node
- It can be used if the tasks of the supervisor program are simpler than those of other tasks.
- If the tasks are independent of each other, they can be scaled well by varying the number of tasks.

# Parallelization strategies/2

## Consecutive

- Each node passes the partially processed data to the next node.
- It can be used if the serial part of the processing is considerably shorter than the parallel part.
- Usually every node runs the same code.
- Particularly suitable for ring topology.

## Parallelization of regions

- Data dependency can be localized into regions.
- It can be used with serial execution time greater than parallel.
- Usually, it needs high volume of communication.
- Most complicated.

# Integral calc. example

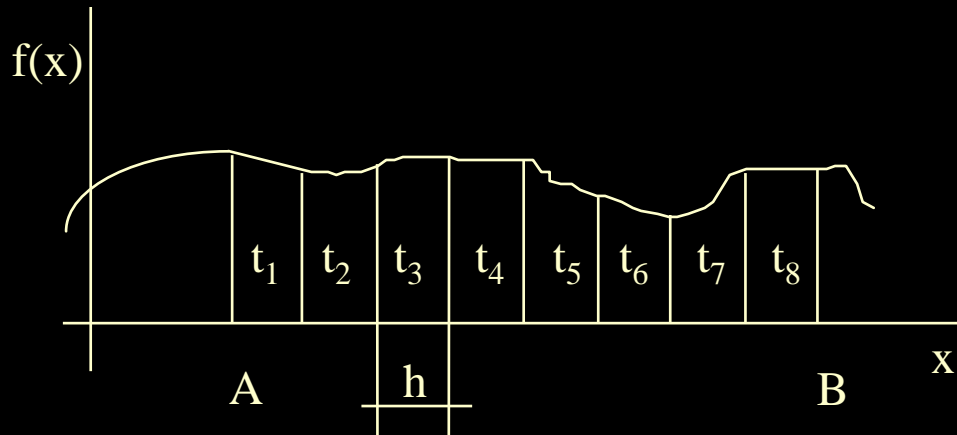
Calculate the

$$\int_A^B f(x) dx$$

integer value by simple numerical approximation (sum of area of rectangles)!

$$\int_A^B f(x) dx = h \cdot \sum_{i=1}^N f\left(A - \frac{h}{2} + i \cdot h\right) \quad \text{with } h = \frac{B - A}{N}$$

In case  $N=8$  eg:



Calculation of area of individual rectangles can be performed independently, in parallel.

# Integral calc. example/2

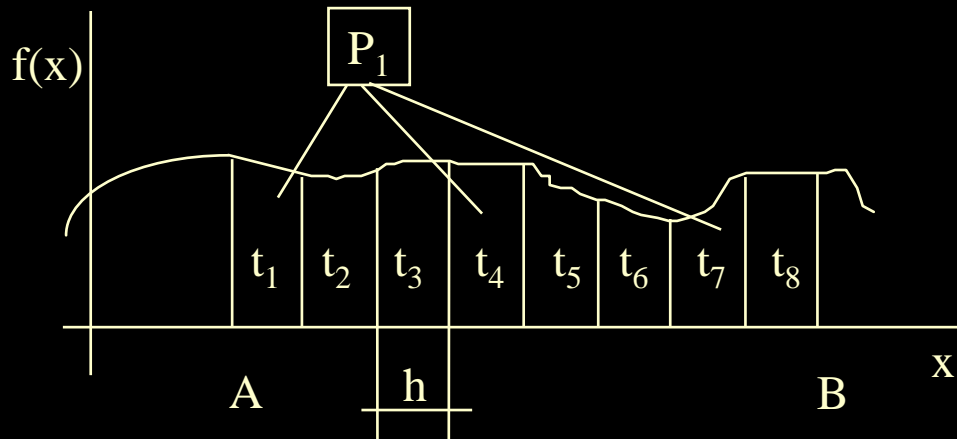
Calculate the

$$\int_A^B f(x) dx$$

integer value by simple numerical approximation (sum of area of rectangles)!

$$\int_A^B f(x) dx = h \cdot \sum_{i=1}^N f\left(A - \frac{h}{2} + i \cdot h\right) \quad \text{where } h = \frac{B - A}{N}$$

In case  $N=8$  eg:



Calculation of area of individual rectangles can be performed independently, in parallel. For example, each task only counts every  $M^{\text{th}}$  area and then sums up the results.

# Integral calc. example/3

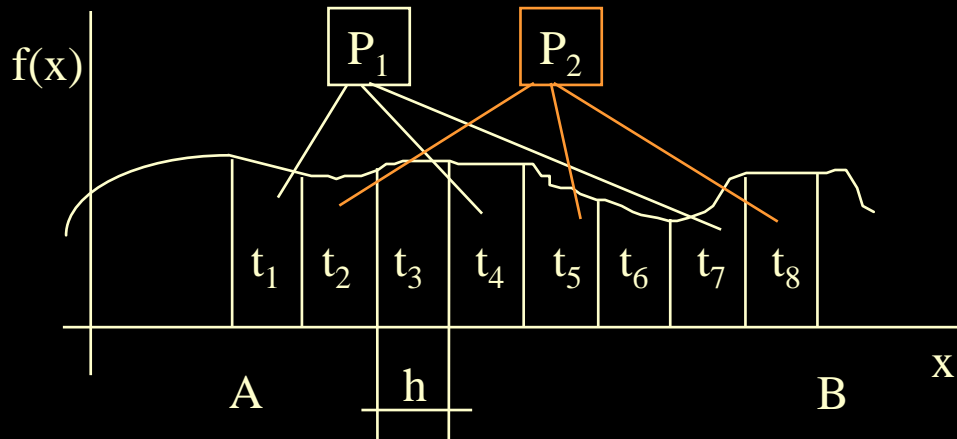
Calculate the

$$\int_A^B f(x) dx$$

integer value by simple numerical approximation (sum of area of rectangles)!

$$\int_A^B f(x) dx = h \cdot \sum_{i=1}^N f\left(A - \frac{h}{2} + i \cdot h\right) \quad \text{where } h = \frac{B - A}{N}$$

In case  $N=8$  eg:



Calculation of area of individual rectangles can be performed independently, in parallel. For example, each task only counts every  $M^{\text{th}}$  area and then sums up the results.

# Integral calc. example/4

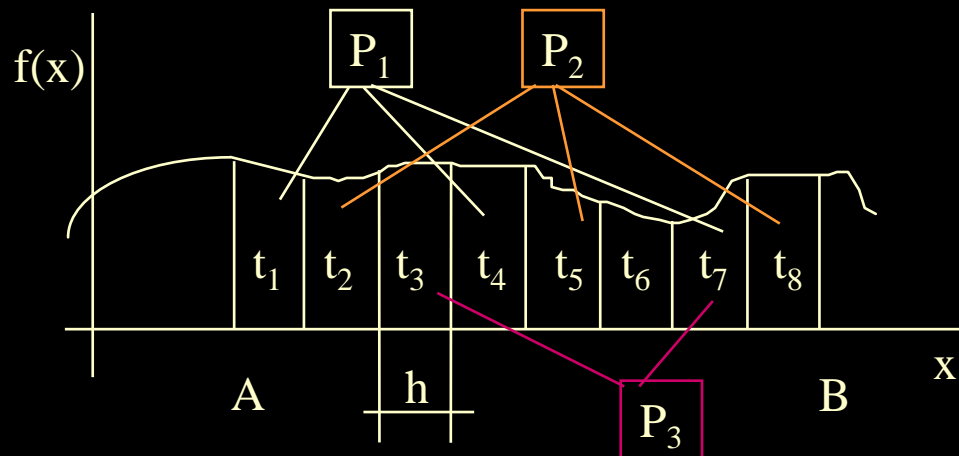
Calculate the

$$\int_A^B f(x) dx$$

integer value by simple numerical approximation (sum of area of rectangles)!

$$\int_A^B f(x) dx = h \cdot \sum_{i=1}^N f\left(A - \frac{h}{2} + i \cdot h\right) \quad \text{where } h = \frac{B - A}{N}$$

In case  $N=8$  eg:



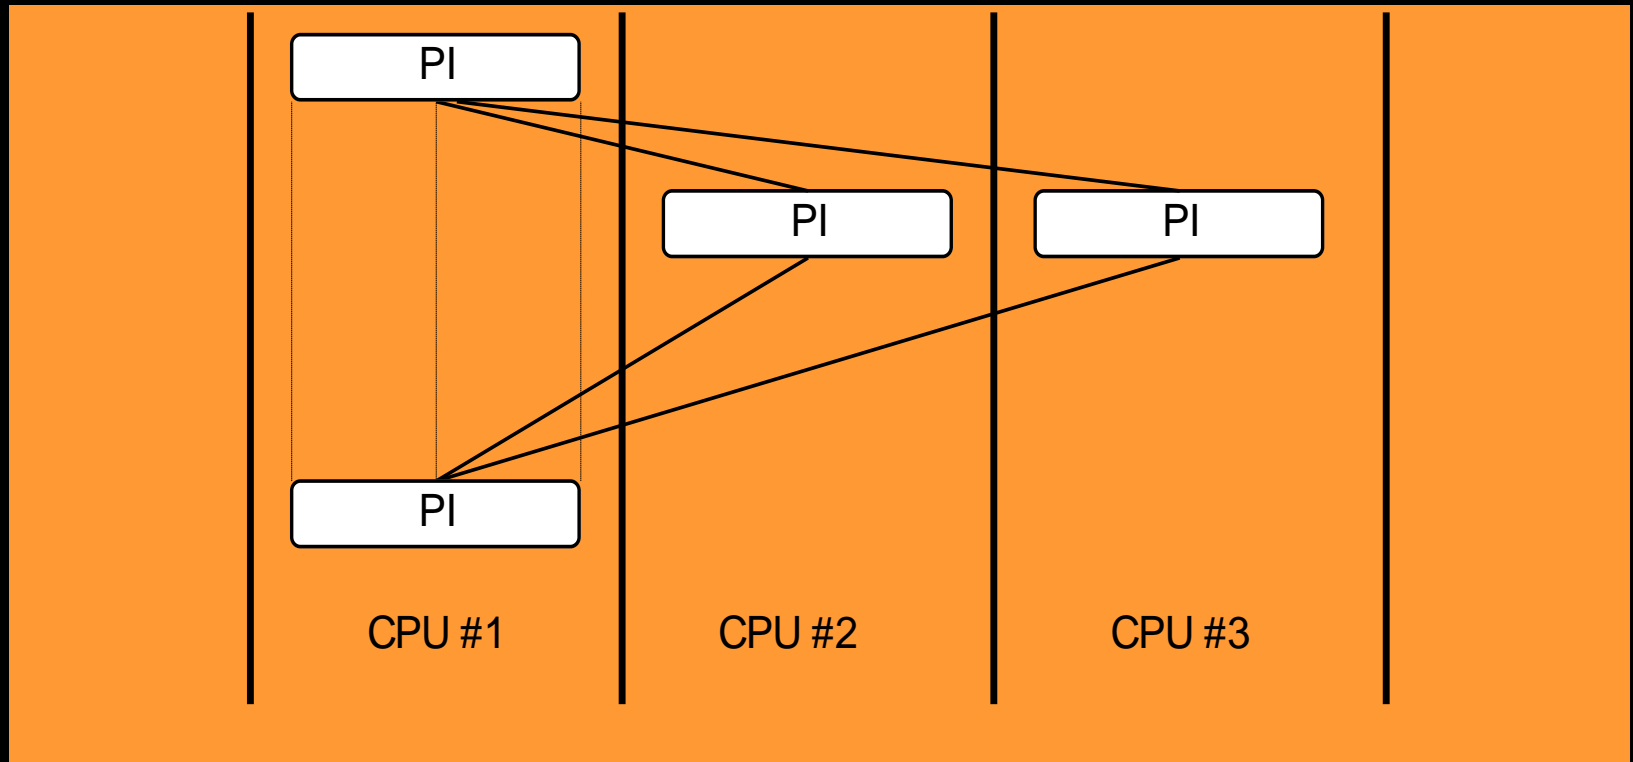
Calculation of area of individual rectangles can be performed independently, in parallel. For example, each task only counts every  $M^{\text{th}}$  area and then sums up the results.



# PI example

Calculate the PI value of a  $\int_0^1 \frac{4}{1+x^2} dx$  integral numerical integration!

Using the above method, we write SPMD program. The first instance of the program prompts the step and starts the other instances. Each instance only counts all  $M^{\text{th}}$  rectangle area, and the results are sent back to the first instance.



# *pi.c program*

```
#include "pvm3.h"          /* PVM 3 include file */
#define f(x) ((float)(4.0/(1.0+x*x))) /* function */
...
// startup function: each instance executes (SPMD)
//   The first one enters to the PVM and starts the other instances
//   Returns: mynum, nprocs, and tids
void startup(int *pmynum, int *pnprocs, int tids[]) {
    int i, mynum, nprocs, info, mytid, numt, parent_tid;
    mytid = pvm_mytid();
    if (mytid < 0) {
        printf("Can't enter to PVM\n");
        exit(0);
    }
    parent_tid = pvm_parent();
```

# *pi.c program /2*

```
if (parent_tid == PvmNoParent) {
    mynum = 0; tids[0] = mytid;
    printf ("How many instance (1-%d)?\n", MAXPROCS);
    scanf("%d", &nprocs);
    numt = pvm_spawn("pi", NULL, PvmTaskDefault, "", nprocs, &tids[1]);
    if (numt != nprocs) { printf ("Error numt != nprocs); exit(0); }
    *pnprocs = nprocs;          /* instance numbers */
    info = pvm_initsend(PvmDataDefault);    /* broadcasts tid info to all */
    info = pvm_pkint(&nprocs, 1, 1);
    info = pvm_pkint(tids, nprocs+1, 1);
    info = pvm_mcast(&tids[1], nprocs, TAG_TIDS);
} else {
    info = pvm_recv(parent_tid, TAG_TIDS);
    info = pvm_upkint(&nprocs, 1, 1);
    info = pvm_upkint(tids, nprocs+1, 1);
    for (i = 1; i <= nprocs; i++) if (mytid == tids[i]) mynum=i;
}
*pmynum = mynum;
}
```

# *pi.c program /3*

```
// solicit function: each instance executes (SPMD)
// Returns the N (numner of rectangles). The first instance asks as a console input
// The others reads from the first instance
void solicit(int *pN, int *pnprocs, int mynum, int tids[]) {
    int info;
    if (mynum == 0) {
        printf("N: (0 = end)\n");
        if (scanf("%d", pN) != 1) *pN = 0;
        info = pvm_initsend(PvmDataDefault);
        info = pvm_pkint(pN, 1, 1);
        info = pvm_pkint(pnprocs, 1, 1);
        info = pvm_mcast(&tids[1], *pnprocs, TAG_N);
    } else {
        info = pvm_recv(tids[0], TAG_N);
        info = pvm_upkint(pN, 1, 1);
        info = pvm_upkint(pnprocs, 1, 1);
    }
}
```

# *pi.c program /4*

```
int main() { // SPMD
    float err, sum, w, x;
    int i, N, M, info, mynum, nprocs, tids[MAXPROCS+1];
    startup(&mynum, &nprocs, tids);
    for (;;) {
        solicit (&N, &nprocs, mynum, tids);
        if (N <= 0) {
            printf("Instance %d exit\n", mynum);
            pvm_exit(); exit(0);
        }

        // Integral approximation
        M = nprocs+1; w = 1.0/(float)N;
        sum = 0.0;
        for (i = mynum+1; i <= N; i += M)
            sum = sum + f(((float)i-0.5)*w);
        sum = sum * w;
```

# *pi.c program /5*

```
// main...
if (mynum == 0) {      /* if it is the first instance */
    printf ("First instance sum = %7.5f\n", sum);
    for (i = 1; i <= nprocs; i++) {
        info = pvm_recv(-1, TAG_SUM);
        info = pvm_upkfloat(&x, 1, 1);
        printf ("First instance got x = %7.5f \n", x); fflush(stdout);
        sum = sum+x;
    }
    printf("sum = %12.8fn", sum); fflush(stdout);
} else {              /* other instances */
    info = pvm_initsend(PvmDataDefault);
    info = pvm_pkfloat(&sum, 1, 1);
    info = pvm_send(tids[0], TAG_SUM);
    printf("A %d. sent his data: %7.2f \n", mynum, sum);
    fflush(stdout);
}
} // main
```

# *Message Passing Interface (MPI)*



- It basically developed with other goals:
  - 1991 -> today
  - Standard, manufacturer-approved, special hw. environment supported by the development environment.
  - long delightful development
  - initially only static process management
  - The communication channels does not set up separately like in the PVM where the daemons must be started before the application.
  - The entire communication layer is linked to the application.