

High Performance Parallel Computing (3.) Threads OpenMP, UPC

Imre Szeberényi, Éva Geist

BME IIT

<szebi@iit.bme.hu>



MŰEGYETEM 1782

Short summary

- Models of parallel computers
 - Flynn's taxonomy
 - idealized parallel computer model
- Programming models
 - Shared memory
 - Distributed shared memory
 - Message passing
- Classes (types) of parallel computers
 - computers with vector processors
 - Symmetric Multiprocessors (SMP)
 - Massively Parallel(MPP)
 - Cluster

Tools (languages)



Tools are not connected to hardware architecture, but their efficiency may differ.

- Linda
- Express
- PVM
- MPI

Design of parallel algorithm

- Computing model (ex. PRAM)
 - algorithm efficiency
- Programming model
 - Processing element (process, task, thread, processor, core)
- Systematic design methods (pl. PACM)
- Design Patterns

Design patterns

- Single Program Multiple Data (SPMD)
 - frequent used method
- Master-Worker (master-slave)
 - fault tolerance, load balancing almost automatic
- Pipeline
 - serial coupled processing element
- Divide and Conquer
- Fork – Join

SPMD pattern

- Same as the SIMD on instruction level
- Mostly results of domain decomposition, but also can be used by control-parallel strategy.
- Supported by the most frequently used tools
- The presented examples (pi, prime) based on SPMD model.
- Sometimes implemented by threads.

Threads



- The memory area of the processes are completely separated.
- The memory area of the threads are common except of stack area.
- In this way the context switch is faster and they can communicate each other easily.
 - `pthread_create()`, `pthread_join()`, `pthread_exit()`
 - `pthread_mutex_...`, `pthread_cond_...`

pthread example

```
#include <stdio.h>
#include <pthread.h>
void* do_loop(void *id)
{
    int i, j;
    float f = 0;
    char me = *(char *)id;
    for (i=0; i<10; i++) {
        for (j=0; j<5000000; j++) f++;
        printf("Thread_%c: %d\n", me, i);
    }
    pthread_exit(NULL);
}
```

thread function

unique data

pthread example/2

```
int main(int argc, char* argv[])
{
    pthread_t thread_a, thread_b;
    char a='a', b='b', c='c';

    pthread_create(&thread_a, NULL,
                  do_loop, &a);
    pthread_create(&thread_b, NULL,
                  do_loop, &b);


    do_loop(&c);
    printf(" can't get here! \n");
    return 0;
}
```

thread handle

unique
parameter

thread function

Race condition problem



Solution: mutual exclusion



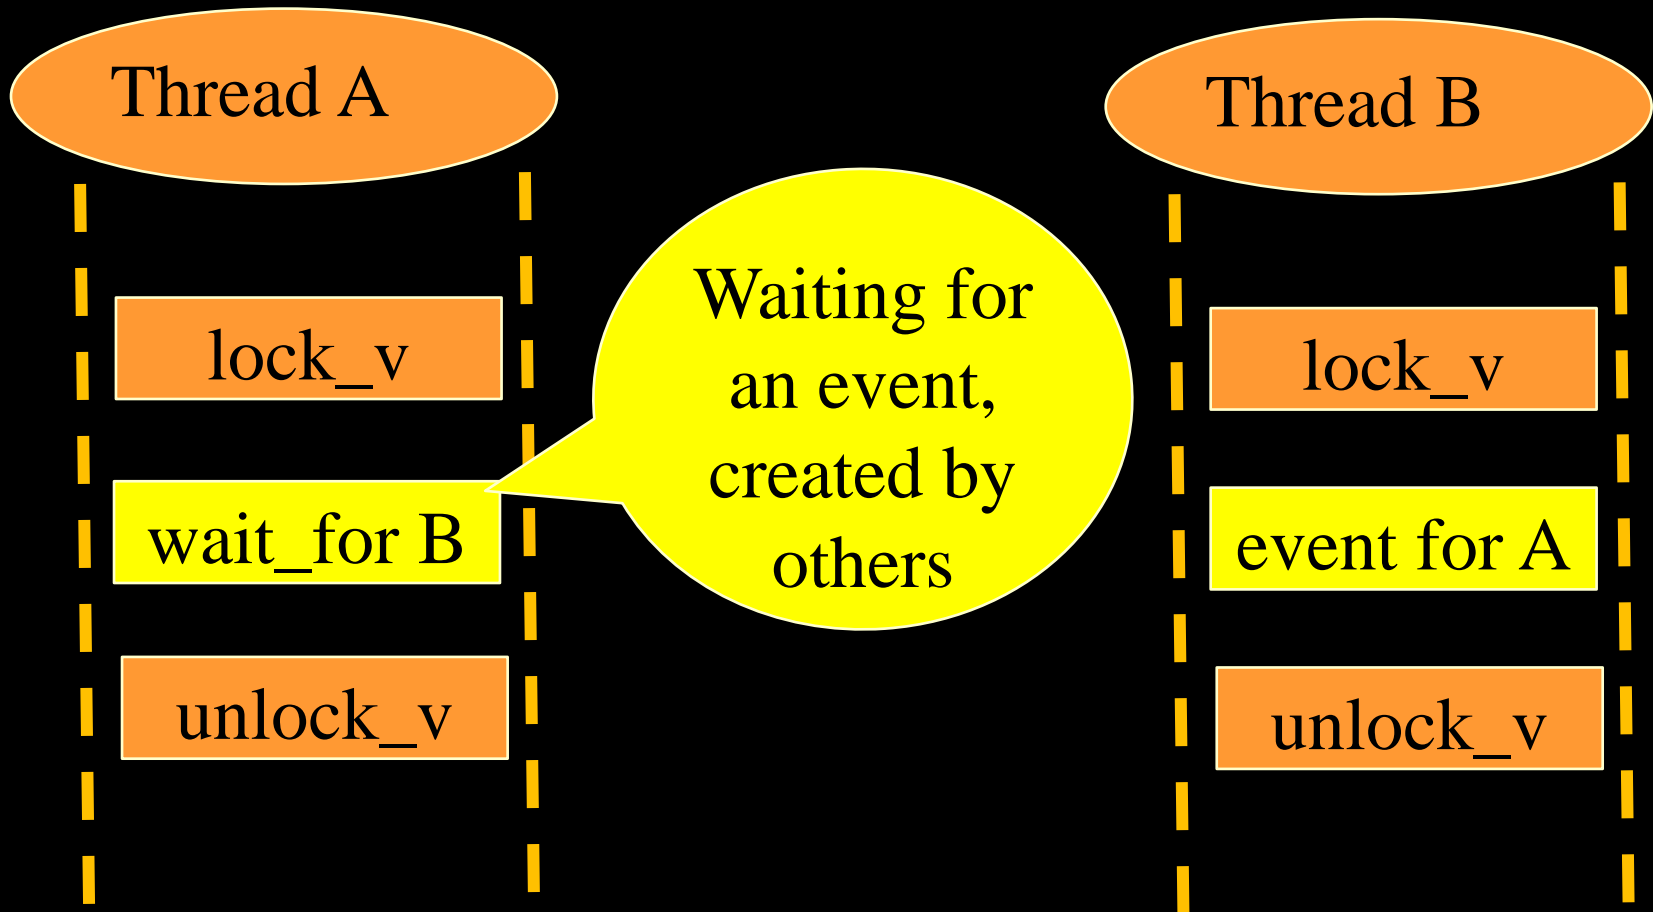
- mutex:
 - pthread_mutex_init
 - pthread_mutex_destroy
 - pthread_mutex_lock
 - pthread_mutex_trylock
 - pthread_mutex_unlock

pthread_mutex



```
....  
pthread_mutex_t mutex_cnt;  
....  
pthread_mutex_init(&mutex_cnt, NULL);  
....  
    pthread_mutex_lock(&mutex_cnt);  
  
    // critical region  
  
    pthread_mutex_unlock(&mutex_cnt);
```

Deadlock problem



Solution: conditional variable

- condition:
 - pthread_cond_init
 - pthread_cond_destroy
 - pthread_cond_wait
 - pthread_cond_timedwait
 - pthread_cond_signal
 - pthread_cond_broadcast

pthread_cond

```
pthread_mutex_t mutex_cnt;
pthread_cond_t cond_cnt;
pthread_mutex_init(&mutex_cnt, NULL);
pthread_cond_init(&cond_cnt, NULL);
..// Thread A
    pthread_mutex_lock(&mutex_cnt);
    ....
    // Critical region; waiting for an event
    pthread_cond_wait(&cond_cnt,
&mutex_cnt);
    ....
    pthread_mutex_unlock(&mutex_cnt);
```

pthread_cond (2)

```
....  
.... // Thread B  
    pthread_mutex_lock(&mutex_cnt);  
        ....  
        // The expected event occurs  
        pthread_cond_signal(&cond_cnt);  
        ....  
    pthread_mutex_unlock(&mutex_cnt);  
....
```


Problems with threads

- Implementation dependent
 - UNIX uses least 2 different implementations
 - Windows (MS) other 2
 - Different language bindings
- No portable solution
- Hard to adapt to the real core numbers
- Not an industry standard

OpenMP




- Language annotation
- The programmer can concentrate to the real problem.
- The parallelization only a possibility
- Shared memory model
- Industrial standard

openmp.org,

computing.llnl.gov/tutorials/openMP,

openmp.org/wp-content/uploads/openmp-examples-4.5.0.pdf

Execution model



Shared memory modell




- The threads are communicate through variables
- The variable sharing are defined on language level
- Race condition problem solution
 - Synchronisation tools
 - Minimising the shared variable usage

Syntax

- `#pragma omp construct [clause [clause] ...]`
- Refers to a block
- OpenMP constructs:
 - Parallel regions
 - Work sharing
 - Data Environment
 - Synchronization
 - Runtime functions/environment variables

Parallel regions

#pragma omp parallel [clause ...] newline
structured_block






- if (scalar_expression)
- private (list)
- shared (list)
- default (shared | none)
- firstprivate (list)
- reduction (operator: list)
- copyin (list)
- num_threads (integer-expression)

Parallel regions example

```
double D[1000] = { 1, 2, 3, 4 };
#pragma omp parallel
{
    int i; double sum = 0;
    for (i=0; i<1000; i++) sum += D[i];
    printf("Thread %d computes %f\n",
          omp_get_thread_num(), sum);
}
// runs as many threads are available
```


The diagram illustrates the memory layout for the provided OpenMP code. A blue box labeled "shared" points to the array `D`, indicating it is shared across all threads. A red box labeled "private" points to the variable `sum`, indicating it is private to each thread.

Work sharing



WS: sections

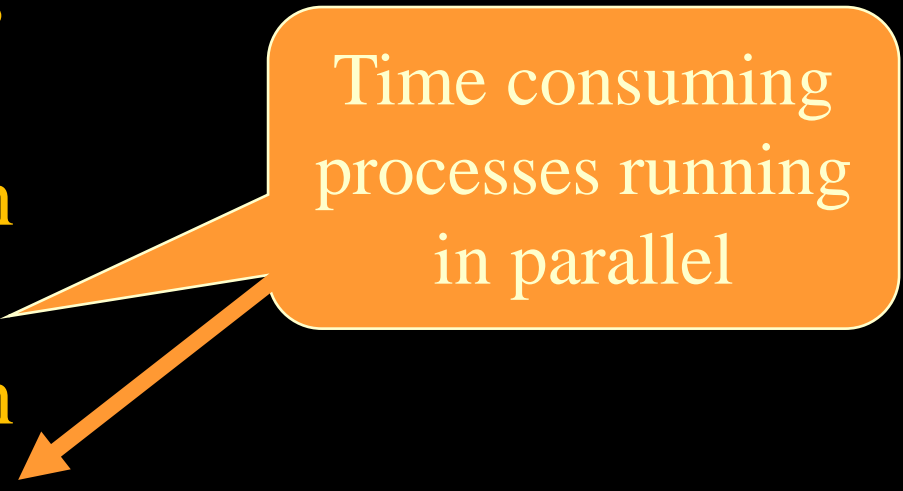
```
#pragma omp sections [clause ...] newline
{
    #pragma omp section newline
        structured_block
    #pragma omp section newline
        structured_block
}
```



```
private (list)
firstprivate (list)
lastprivate (list)
reduction (operator: list)
nowait
```

Sections example

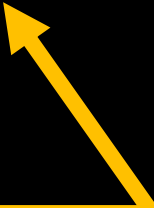
```
#pragma omp sections
{
    #pragma omp section
    a = computation_1();
    #pragma omp section
    b = computation_2();
}
c = a + b;
```



Time consuming
processes running
in parallel

WS: single

#pragma omp single [clause ...] newline
structured_block



private (list)
firstprivate (list)
nowait

Cycle parallelization by „hand”

```
for (int i=0; i<N; i++) { a[i]=b[i]+c[i]; }
```

```
#pragma omp parallel
{
    int id = omp_get_thread_num();
    int nThr = omp_get_num_threads();
    int istart = id*N/nThr, iend = (id+1)*N/nThr;
    for (int i=istart; i<iend; i++) { a[i]=b[i]+c[i]; }
}
```

Automatic parallelization

```
#pragma omp parallel
#pragma omp for schedule(static)
{
    for (int i=0; i<N; i++) { a[i]=b[i]+c[i]; }
}
```

Applicable on simple for only:

- only one int loop variable,
- cmp. op: <, >, <=, >
- incrementing/decrementing: ++, --, assignmet op.
- cycle independent inic, cond, and increment

WS: for

#pragma omp for [clause ...] newline
for_loop



Schedule (type [,chunk])

ordered

private (list)

firstprivate (list)

lastprivate (list)

shared (list)

reduction (operator: list)

collapse

nowait

Scheduling

- `schedule(static [, chunksize])`
 - static divisions
 - Default: same sizes
 - Round-robin (more slices than threads)
- `schedule(dynamic [, chunksize])`
 - dynamic slicing
 - Default chunksize = 1
- `schedule(guided [, chunksize])`
 - Dynamic, exponential

Granularity

- `#pragma omp parallel if (expression)`
 - condition for parallelization
- `#pragma omp num_threads (expression)`
 - the number of threads can be modified

Data Environment

- Shared variables
 - global variables

```
int sum = 0;
#pragma omp parallel for
for (int i=0; i<N; i++) sum += i;
```
- Private
 - automatic variables in parallel blokk
 - automatic variables in functions
 - explicit private declaration

Variable types

- **private:**
 - private instance without initialization
 - same as it would be in { }
- **firstprivate:**
 - private initialized from the variable
- **lastprivate:**
 - value of last loop copied back
- **threadprivate:**
 - persistent private

Variable examples

```
int i;  
#pragma omp parallel for private(i)  
for (i=0; i<n; i++) { ... }
```

```
int idx=1;  
int x = 10;  
#pragma omp parallel for  
#pragma omp firstprivate(x) lastprivate(idx)  
for (i=0; i<n; i++) {  
    if (data[i]==x) idx = i;  
}
```

Variable examples /2

```
int data[100];  
#pragma omp threadprivate(data)  
...  
#pragma omp parallel for copyin(data)  
for (int i=0; i<n; i++)  
    data[i]++;
```

Reduction

```
int sum = 0;
#pragma omp parallel for reduction(+: sum)
for (int i = 0; i < N; i++)
    sum = sum + a[i] * b[i];
```

- Only skalar variable
- each thread has a temporary
- the form is: $x \text{ op expr}$
- $x++$, $++x$, $x--$, $--x$,
- op could not been overloaded

Synchronization

- Single/Master execution
 - #pragma omp single
 - #pragma omp master
- Critical sections, Atomic updates
 - #pragma omp critical [name]
 - #pragma omp atomic
 - update_statement
 - only scalar

Synchronization /2

- Ordered

`#pragma omp ordered`

```
int vec[100];  
#pragma omp parallel for ordered  
for (int i=0; i<100; i++) {  
    vec[i] = 2 * vec[i];  
    #pragma omp ordered  
    printf("vec[i] = %d\n", vec[i]);  
}
```

Synchronization /3

- Barriers

`#pragma omp barrier`

- Nowait

`#pragma omp sections nowait`

- Flush

`#pragma omp flush (list)`

- Reduction

Summary of directives

Clause	Directive					
	PARALLEL	DO/for	SECTIONS	SINGLE	PARALLEL DO/for	PARALLEL SECTIONS
IF	●				●	●
PRIVATE	●	●	●	●	●	●
SHARED	●	●			●	●
DEFAULT	●				●	●
FIRSTPRIVATE	●	●	●	●	●	●
LASTPRIVATE		●	●		●	●
REDUCTION	●	●	●		●	●
COPYIN	●				●	●
COPYPRIVATE				●		
SCHEDULE		●			●	
ORDERED		●			●	
NOWAIT		●	●	●		

Control functions

- `omp_set_dynamic(int)/ omp_get_dynamic()`
- `omp_set_num_threads(int)/ omp_get_num_threads()`
 - `OMP_NUM_THREADS` env.
- `omp_get_num_procs()`
- `omp_get_thread_num()`
- `omp_set_nested(int)/omp_get_nested()`
- `omp_in_parallel()`
- `omp_get_wtime()`
- `omp_init_lock()`, `omp_destroy_lock()`,
- `omp_set_lock()`, `omp_unset_lock()`,
- `omp_test_lock()`

Unified Parallel C (UPC)

- extension of C99
- Shared memory model
- supports the NUMA
- development started in 1999.
- most supercomputers support it.

upc-lang.org, upc.lbl.gov,
upc.gwu.edu/tutorials/UPC-SC05.pdf


Main features of UPC

- Explicit parallel model
 - Independent threads (thread != OS thread)
 - Each thread complies to the C standard (many main).
- Partitioned Global Address Space(PGAS)
 - shared and private data
 - scalars are in the p0, but arrays are distributed
- Synchronization primitives
- Runtime library
 - scatter, gather, reductions, exchange, ..

UPC keywords

- MYTHREAD, THREADS
- shared, relaxed, strict,
- upc_barrier, upc_notify, upc_wait
- upc_forall,
- upc_blocksizeof, upc_elementsizeof,
upc_localsizeof,
UPC_MAXB_LOCK_SIZE
- upc_fence,

Different memory models



Partitioned Global Address Sp.

- The shared arrays are distributed in the memory associated to the processing elements (affinity)
- The threads can work more efficiently on the „nearest” data.

```
shared int a[100]; // cyclic distribution
upc_forall(int i = 0; i < 100; i++; &a[i]) ...
// The given cycle run on the nearest data
```

Affinity and block size

```
shared [block_size] type array[N];  
affinity = (i/block_size)%THREADS
```

```
shared [3] int A[4][THREADS];
```

```
supposed: THREADS = 4
```

Thread 0

A[0][0]
A[0][1]
A[0][2]
A[3][0]
A[3][1]
A[3][2]

Thread 1

A[0][3]
A[1][0]
A[1][1]
A[3][3]

Thread 2

A[1][2]
A[1][3]
A[2][0]

Thread 3

A[2][1]
A[2][2]
A[2][3]

example: PI

```
shared long hits[THREADS]; // counters
```

```
....
```

```
upc_forall (i=0; i < my_trials; i++; continue)
```

```
    hits[MYTHREAD] += isInside();
```

```
upc_barrier;           // wait for all
```

```
if (MYTHREAD == 0) // summ of the local cnts
```

```
    for (i = 1; i < THREADS; i++) hits[0] += hits[i];
```

Examples



para.iit.bem.hu/~szebi/para/threads

Para.iit.bem.hu/~szebi/para/OpenMP

Para.iit.bem.hu/~szebi/para/UPC

**# For using UPC compiler the environment
should be set first:**

Module load upc